Understanding your AWS Account

1. Which of the following EC2 services can be used without charge under the Free Tier?

A. Any single EC2 instance type as long as it runs for less than one hour per day

B. Any single EC2 instance type as long as it runs for less than 75 hours per month

C. A single t2.micro EC2 instance type instance for 750 hours per month

D. t2.micro EC2 instance type instances for a total of 750 hours per month

2. You want to experiment with deploying a web server on an EC2 instance. Which two of the following resources can you include to make that work while remaining within the Free Tier? (Select TWO.)

A. A 5 GB bucket on S3

B. A t2.micro instance type EC2 instance

C. A 30 GB solid-state Elastic Block Store (EBS) drive

D. Two 20 GB solid-state Elastic Block Store (EBS) drives

3. Which of the following usage will always be cost-free even after your account's Free Tier has expired? (Select TWO.)

A. One million API calls/month on Amazon API Gateway

B. 10 GB of data retrievals from Amazon Glacier per month

C. 500 MB/month of free storage on the Amazon Elastic Container Registry (ECR)

D. 10 custom monitoring metrics and 10 alarms on Amazon CloudWatch

4. Which of the following tools are available to ensure you won't accidentally run past your Free Tier limit and incur unwanted costs? (Select TWO.)

A. Automated email alerts when activity approaches the Free Tier limits

B. The Top Free Tier Services by Usage section on the Billing & Cost Management Dashboard

C. Billing & Cost Management section on the Top Free Tier Services Dashboard

D. The Billing Preferences Dashboard

5. Which of the following is likely to be an accurate source of AWS pricing information?

A. Wikipedia pages relating to a particular service

B. The AWS Command Line Interface (AWS CLI)

C. **AWS online documentation relating to a particular service**

D. The AWS Total Cost of Ownership Calculator

6. Which of the following will probably not affect the pricing for an AWS service?

A. Requests for raising the available service limit

B. AWS Region

C. The volume of data saved to an S3 bucket

D. The volume of data egress from an Amazon Glacier vault

7. Which of the following is a limitation of the AWS Simple Monthly Calculator?

A. You can calculate resource use for only one service at a time.

B. Not all AWS services are included.

C. The pricing is seldom updated and doesn't accurately reflect current pricing.

D. You're not able to specify specific configuration parameters.

8. Which of the following Simple Monthly Calculator selections will likely have an impact on most other configuration choices on the page? (Select TWO.)

A. Calculate By Month Or Year

B. Include Multiple Organizations

C. Free Usage Tier

D. Choose Region

9. Which of the following is not an included parameter in the AWS Total Cost of Ownership Calculator?

A. The tax implications of a cloud deployment

B. Labor costs of an on-premises deployment

C. Networking costs of an on-premises deployment

D. Electricity costs of an on-premises deployment

10. Which of the following AWS Total Cost of Ownership Calculator parameters is likely to have the greatest impact on cost?

A. Currency

B. AWS Region

C. Guest OS

D. Number of servers

11. Which of the following AWS documentation URLs points to the page containing an up-to-date list of service limits?

A. https://docs.aws.amazon.com/general/latest/gr/limits.html

B. https://docs.aws.amazon.com/general/latest/gr/aws_service_limits.html

C. https://aws.amazon.com/general/latest/gr/aws_service_limits.html

D. https://docs.aws.amazon.com/latest/gr/aws_service_limits.html

12. Which of the following best describes one possible reason for AWS service limits?

A. To prevent individual customers from accidentally launching a crippling level of resource consumption

B. To more equally distribute available resources between customers from different regions

C. To allow customers to more gradually increase their deployments

D. Because there are logical limits to the ability of AWS resources to scale upward

13. Is it always possible to request service limit increases from AWS?

A. Yes. All service limits can be increased.

B. No. A limit can never be increased.

C. Service limits are defaults. They can be increased or decreased on demand.

D. No. Some service limits are hard.

14. Which is the best place to get a quick summary of this month's spend for your account?

A. Budgets

B. Cost Explorer

C. Cost and usage reports

D. Billing & Cost Management Dashboard

15. What is the main goal for creating a Usage budget type (in AWS Budgets)?

A. To correlate usage per unit cost to understand your account cost efficiency

B. To track the status of any active reserved instances on your account

C. To track particular categories of resource consumption

D. To monitor costs being incurred against your account

16. Which of the following is not a setting you can configure in a Cost budget?

A. Period (monthly, quarterly, etc.)

B. Instance type

C. Start and stop dates

D. Owner (username of resource owner)

17. What is the main difference between the goals of Cost Explorer and of cost and usage reports?

A. Cost Explorer displays visualizations of high-level historical and current account costs, while cost and usage reports generate granular usage reports in CSV format.

B. Cost and usage reports display visualizations of high-level historical and current account costs, while Cost Explorer generates granular usage reports in CSV format.

C. Cost Explorer lets you set alerts that are triggered by billing events, while cost and usage reports help you visualize system events.

D. Cost and usage reports are meant to alert you to malicious intrusions, while Cost Explorer displays visualizations of high-level historical and current account costs.

18. What is the purpose of cost allocation tags?

A. To associate spend limits to automatically trigger resource shutdowns when necessary

B. To help you identify the purpose and owner of a particular running resource to better understand and control deployments

C. To help you identify resources for the purpose of tracking your account spending

D. To visually associate account events with billing periods

19. Which of the following scenarios would be a good use case for AWS Organizations? (Select TWO.)

A. A single company with multiple AWS accounts that wants a single place to administrate everything

B. An organization that provides AWS access to large teams of its developers and admins

C. A company that's integrated some operations with an upstream vendor

D. A company with two distinct operational units, each with its own accounting system and AWS account

20. Which of these tools lets you design graphs within the browser interface to track your account spending?

A. Budgets

B. Cost Explorer

C. Reports

D. Consolidating Billing

Core Compute Services

1. What is the function of an EC2 AMI?

A. To define the hardware profile used by an EC2 instance

B. To serve as an instance storage volume for high-volume data processing operations

C. To serve as a source image from which an instance's primary storage volume is built

D. To define the way data streams are managed by EC2 instances

2. Where can you find a wide range of verified AMIs from both AWS and third-party vendors?

A. AWS Marketplace

B. Quick Start

C. Community AMIs

D. My AMIs

3. Which of the following could be included in an EC2 AMI? (Select TWO.)

A. A networking configuration

B. A software application stack

C. An operating system

D. An instance type definition

4. Which of the following are EC2 instance type families? (Select TWO.)

A. c5d.18xlarge

B. Compute optimized

C. t2.micro

D. Accelerated computing

5. When describing EC2 instance types, what is the role played by the vCPU metric?

A. vCPUs represent an instance's potential resilience against external network demands.

B. vCPUs represent an instance type's system memory compared to the class of memory modules on a physical machine.

C. vCPUs represent an AMI's processing power compared to the number of processors on a physical machine.

D. vCPUs represent an instance type's compute power compared to the number of processors on a physical machine.

6. Which of the following describes an EC2 dedicated instance?

A. An EC2 instance running on a physical host reserved for the exclusive use of a single AWS account
B. An EC2 instance running on a physical host reserved for and controlled by a single AWS account

C. An EC2 AMI that can be launched only on an instance within a single AWS account

D. An EC2 instance optimized for a particular compute role

7. Which of the following describes an EBS volume?

A. A software stack archive packaged to make it easy to copy and deploy to an EC2 instance

B. A virtualized partition of a physical storage drive that's directly connected to the EC2 instance it's associated with

C. A virtualized partition of a physical storage drive that's not directly connected to the EC2 instance it's associated with

D. A storage volume that's encrypted for greater security

8. Why might you want to use an instance store volume with your EC2 instance rather than a volume from the more common service? (Select TWO.)

A. Instance store volumes can be encrypted.

B. Instance store volumes, data will survive an instance shutdown.

C. Instance store volumes provide faster data read/write performance.

D. Instance store volumes are connected directly to your EC2 instance.

9. Your web application experiences periodic spikes in demand that require the provisioning of extra instances. Which of the following pricing models would make the most sense for those extra instances?

A. Spot

B. On-demand

C. Reserved

D. Dedicated

10. Your web application experiences periodic spikes in demand that require the provisioning of extra instances. Which of the following pricing models would make the most sense for the “base” instances that will run constantly?

A. Spot

B. On-demand

C. Spot fleet

D. Reserved

11. Which of the following best describes what happens when you purchase an EC2 reserved instance?

A. Charges for any instances you run matching the reserved instance type will be covered by the reservation.

B. Capacity matching the reserved definition will be guaranteed to be available whenever you request it.

C. Your account will immediately and automatically be billed for the full reservation amount.

D. An EC2 instance matching your reservation will automatically be launched in the selected AWS Region.

12. Which of the following use cases are good candidates for spot instances? (Select TWO.)

A. Big data processing workloads

B. Ecommerce websites

C. Continuous integration development environments

D. Long-term, highly available, content-rich websites

13. Which AWS services simplify the process of bringing web applications to deployment? (Select TWO.)

A. Elastic Block Store

B. Elastic Compute Cloud

C. Elastic Beanstalk

D. Lightsail

14. Which of the following services bills at a flat rate regardless of how it's consumed?

A. Lightsail

B. Elastic Beanstalk

C. Elastic Compute Cloud

D. Relational Database Service

15. Which of these stacks are available from Lightsail blueprints? (Select TWO.)

A. Ubuntu

B. Gitlab

C. WordPress

D. LAMP

16. Which of these AWS services use primarily EC2 resources under the hood? (Select TWO.)

A. Elastic Block Store

B. Lightsail

C. Elastic Beanstalk

D. Relational Database Service

17. Which of the following AWS services are designed to let you deploy Docker containers? (Select TWO.)

A. Elastic Container Service

B. Lightsail

C. Elastic Beanstalk

D. Elastic Compute Cloud

18. Which of the following use container technologies? (Select TWO.)

A. Docker

B. Kubernetes

C. Lambda

D. Lightsail

19. What role can the Python programming language play in AWS Lambda?

A. Python cannot be used for Lambda.

B. It is the primary language for API calls to administrate Lambda remotely.

C. It is used as the underlying code driving the service.

D. It can be set as the runtime environment for a function.

20. What is the maximum time a Lambda function may run before timing out?

A. 15 minutes

B. 5 minutes

C. 1 minute

D. 1 hour